

DJNW Metadata Audit

How reliable is Dow Jones's article tagging?

April 17, 2026

2026-04-17

DJNW Metadata Audit

DJNW Metadata Audit

How reliable is Dow Jones's article tagging?

April 17, 2026

Motivation: DJNW metadata is unreliable

DJNW tagging scheme: full census

2026-04-17

Motivation: DJNW metadata is unreliable

DJNW tagging scheme: full census

Motivation: DJNW metadata is unreliable

2026-04-17

DJNW Metadata Audit
└ Motivation: DJNW metadata is unreliable

Motivation: DJNW metadata is unreliable

Case study: crude oil tags in DJNW

Dow Jones tags each article with metadata codes. We investigated how well these tags identify crude-oil-relevant articles across the full 1979–2025 archive (31.7M articles, 555 months).

Scanning all five taxonomy CSVs, we identified **16 content codes** plus **3 oil/energy wire codes** that could identify crude-oil articles:

Code	Description	Code	Description
<i>Content codes (what the article is about)</i>			
N/PET	Crude Oil & Petroleum Products	I/OIL	Major Oil & Gas
N/CMKT	Crude Spot Market Commentary	I/OIS	Oil Extraction
N/PRD	Oil & Natural Gas Production	I/FSL	Fossil Fuels
N/OPC	OPEC	M/ENE	Energy market
N/REF	Refinery Outages	S/OIL	Oil
N/RMKT	Refined Products Spot Market Commentary	S/API	API Weekly Statistics
N/RPRI	Refined Products Spot Market Prices	S/DOE	DOE Weekly Data
N/ENY	Energy	N/EGY	Energy Commentary
<i>Wire codes (which DJ service carried the article)</i>			
N/IPR	International Petroleum Report	N/DOI	DJ Oil & Gas Service
N/NRG	Dow Jones Energy Service		

Note: “WTI” and “Brent” are *not* DJNW codes — they are words in headlines. DJNW uses N/PET for both.

DJNW Metadata Audit


└ Motivation: DJNW metadata is unreliable

└ Case study: crude oil tags in DJNW

2026-04-17

Case study: crude oil tags in DJNW	
Dow Jones tags each article with metadata codes. We investigated how well these tags identify crude-oil-relevant articles across the full 1979–2025 archive (31.7M articles, 555 months).	
Scanning all five taxonomy CSVs, we identified 16 content codes plus 3 oil/energy wire codes that could identify crude-oil articles.	
Code	Description
Content codes (what the article is about)	
N/PET	Crude Oil & Petroleum Products
N/CMKT	Crude Spot Market Commentary
N/PRD	Oil & Natural Gas Production
N/OPC	OPEC
N/REF	Refinery Outages
N/RMKT	Refined Products Spot Market Commentary
N/RPRI	Refined Products Spot Market Prices
N/ENY	Energy
Wire codes (which DJ service carried the article)	
N/IPR	International Petroleum Report
N/NRG	Dow Jones Energy Service
Note: “WTI” and “Brent” are not DJNW codes — they are words in headlines. DJNW uses N/PET for both.	

Content codes: absolute usage over time




Log scale, monthly. Most content codes do not appear until the mid-1990s.


2026-04-17

DJNW Metadata Audit

- └ Motivation: DJNW metadata is unreliable
- └ Content codes: absolute usage over time




Wire codes: absolute usage over time




Log scale, monthly. Wire codes saturate at $\sim 100k$ articles/month from 1979 onwards (they were stamped on everything), then drop sharply at Nov 1996 and stabilize at a lower volume.

2026-04-17

- DJNW Metadata Audit
 - Motivation: DJNW metadata is unreliable
 - Wire codes: absolute usage over time



Content codes: share of articles carrying each tag




Baseline is sparse — no code exceeds ~8% even at peaks.

2026-04-17


DJNW Metadata Audit

└ Motivation: DJNW metadata is unreliable

└ Content codes: share of articles carrying each tag



Wire codes: share of articles carrying each tag




Wire codes: 100% spam pre-Nov 1996, brief regressions in Mar/Apr/Oct–Dec 2000, then selective (~10–20%) from 2001. Post-2000, N/IPR is a useful weak positive signal.

2026-04-17

DJNW Metadata Audit

└ Motivation: DJNW metadata is unreliable

└ Wire codes: share of articles carrying each tag



Recall failure 1: direct oil reporting without oil tags

Articles where crude oil *is the primary subject* yet carry **none** of the 19 oil codes:

- “Economic Outlook 99: Saudi Arabia Hurt By Low Oil Prices” (1998-12-14)
- “Saudi Oil Price Cut Dings Energy ETFs” (2014-11-04, Barron's)
- “Venezuela Serves Up a Culture Shock for Visiting Saudi Oil Minister” (2014-11-07)
- “Russia to Pump Up Oil Exports to Asia” (2014-12-04)
- “Abu Dhabi’s Mubadala Wealth Fund Swings to Loss on Low Oil Prices” (2016-09-08)
- “Nymex Crude Biased Down Near-Term; \$78.08 Support” (2014-11-03)
- “Exxon to Shut Down Oil Production in Russia After Ukraine Attack” (2022-03-01)

These are exactly the articles a macro researcher searching for OPEC, Saudi, or petro-state policy *should* find by tag. All missed.

2026-04-17

DJNW Metadata Audit
└ Motivation: DJNW metadata is unreliable

└ Recall failure 1: direct oil reporting without oil tags

Verification (all in ../rc2573/output/cleaned/v2/articles/):

19981214001135 in 1998-12a_clean.jsonl
20141104009555, 20141107013622 in 2014-11a_clean.jsonl
20141204001090 in 2014-12a_clean.jsonl
20160908005314 in 2016-09a_clean.jsonl
20141103015258 in 2014-11a_clean.jsonl
20220301015891 in 2022-03a_clean.jsonl

All verified against the 19-code set (16 content + N/IPR, N/NRG, N/DOI).

Full results: results/missing_oil_tags_v3.txt (930 DIRECT cases across 20 event months).

Recall failure 1: direct oil reporting without oil tags

Articles where crude oil is the primary subject yet carry none of the 19 oil codes:

- “Economic Outlook 99: Saudi Arabia Hurt By Low Oil Prices” (1998-12-14)
- “Saudi Oil Price Cut Dings Energy ETFs” (2014-11-04, Barron's)
- “Venezuela Serves Up a Culture Shock for Visiting Saudi Oil Minister” (2014-11-07)
- “Russia to Pump Up Oil Exports to Asia” (2014-12-04)
- “Abu Dhabi’s Mubadala Wealth Fund Swings to Loss on Low Oil Prices” (2016-09-08)
- “Nymex Crude Biased Down Near-Term; \$78.08 Support” (2014-11-03)
- “Exxon to Shut Down Oil Production in Russia After Ukraine Attack” (2022-03-01)

These are exactly the articles a macro researcher searching for OPEC, Saudi, or petro-state policy should find by tag. All missed.

Recall failure 2: cross-asset articles reacting to oil

Articles about FX, equities, or bonds *reacting to oil* also frequently lack oil codes — even though the article's thesis is literally “asset *X* moves because of oil”:

- “Canadian Dollar Ends Sharply Down on Lower Oil Prices, Dovish Poloz” (2014-11-03)
- “Ruble Hits New Low as Oil Prices Drop” (2014-12-01)
- “NZ Shares Fall as Saudi Attack Sends Oil Price Gushing” (2019-09-15, post-Abqaiq)
- “Shares in European Airlines Fall as Oil Price Jumps” (2020-01-03, post-Soleimani)
- “Falling Oil Prices Help Cruise, Airline Stocks” (2023-10-04, Israel-Gaza)
- “Stock Futures Fall as Oil Rises Back Above \$100” (2022-03-01)

Arguably gray-zone — DJ might reasonably not tag these as “oil” (the subject is the reacting asset). But for research on oil’s macroeconomic effects, these are exactly the articles we’d want surfaced.

DJNW Metadata Audit
└ Motivation: DJNW metadata is unreliable

└ Recall failure 2: cross-asset articles reacting to oil

Verification (all in ../rc2573/output/cleaned/v2/articles/):

20141103012721 in 2014-11a_clean.jsonl
20141201007859 in 2014-12a_clean.jsonl
20190915001150 in 2019-09b_clean.jsonl
20200103001619 in 2020-01a_clean.jsonl
20231004005536 in 2023-10a_clean.jsonl
20220301002641 in 2022-03a_clean.jsonl

Full results: results/missing_oil_tags.v3.txt (1,591 CROSS_ASSET cases).

Articles about FX, equities, or bonds reacting to oil also frequently lack oil codes — even though the article's thesis is literally “asset *X* moves because of oil”:

- “Canadian Dollar Ends Sharply Down on Lower Oil Prices, Dovish Poloz” (2014-11-03)
- “Ruble Hits New Low as Oil Prices Drop” (2014-12-01)
- “NZ Shares Fall as Saudi Attack Sends Oil Price Gushing” (2019-09-15, post-Abqaiq)
- “Shares in European Airlines Fall as Oil Price Jumps” (2020-01-03, post-Soleimani)
- “Falling Oil Prices Help Cruise, Airline Stocks” (2023-10-04, Israel-Gaza)
- “Stock Futures Fall as Oil Rises Back Above \$100” (2022-03-01)

Arguably gray zone — DJ might reasonably not tag these as “oil” (the subject is the reacting asset). But for research on oil’s macroeconomic effects, these are exactly the articles we’d want surfaced.

The wire-code problem: pre-1996 spam and the 1996 transition

The N/ prefix is **overloaded** — it mixes content tags with wire codes (which DJ service carried the article). In 1979–1996, every article was stamped with ~19 identical wire codes including N/IPR, N/NRG, N/DOI. Non-energy articles carrying N/IPR:

Date	Headline (all have N/IPR)
1985-06	<i>"IBM Announces New Modular Processor for Finance Industry"</i>
1990-03	<i>"NV Philips Year Net Rose 30% After Gains"</i>
1993-06	<i>"Richmond Fed's Broadbus: US Econ Poised For Moderate Growth"</i>
1993-06	<i>"UK May Purchasing Mgrs Index Fell To 53.0%"</i>

In Nov 1996, N/IPR coverage drops from 100% → 0.4% in one month (the same step appears across N/NRG, N/DOI — a tagging-system change, *inferred*). Post-1996 it stabilizes at 10–20% and becomes a useful weak positive signal for oil content.

In practice: include N/IPR, N/NRG, N/DOI post-2000 as weak positive signals. Pre-Nov 1996 and Mar/Apr/Oct–Dec 2000 are anomalous and should be masked.

DJNW Metadata Audit

└ Motivation: DJNW metadata is unreliable

└ The wire-code problem: pre-1996 spam and the 1996 transition

Verification references (all in ../rc2573/output/cleaned/v2/articles/):

1985-06_clean.jsonl: 4,038/4,038 articles carry the same 19 wire codes.

19850603000003 IBM 19900301000002 NV Philips

19930601000002 Richmond Fed 19930601000004 UK May PMI

Share of articles carrying N/IPR: 1995-12 = 100.0%, 1996-05 = 100.0%, 1996-11 = 0.4%, 1997-01 = 0.3%, 2022-03 = 14.8%.

The wire-code problem: pre-1996 spam and the 1996 transition

The N/ prefix is **overloaded** — it mixes content tags with wire codes (which DJ service carried the article). In 1979–1996, every article was stamped with ~19 identical wire codes including N/IPR, N/NRG, N/DOI. Non-energy articles carrying N/IPR:

Date	Headline (all have N/IPR)
1985-06	<i>"IBM Announces New Modular Processor for Finance Industry"</i>
1990-03	<i>"NV Philips Year Net Rose 30% After Gains"</i>
1993-06	<i>"Richmond Fed's Broadbus: US Econ Poised For Moderate Growth"</i>
1993-06	<i>"UK May Purchasing Mgrs Index Fell To 53.0%"</i>

In Nov 1996, N/IPR coverage drops from 100% → 0.4% in one month (the same step appears across N/NRG, N/DOI — a tagging-system change, *inferred*). Post-1996 it stabilizes at 10–20%, and becomes a useful weak positive signal for oil content.

In practice: include N/IPR, N/NRG, N/DOI post-2000 as weak positive signals. Pre-Nov 1996 and Mar/Apr/Oct–Dec 2000 are anomalous and should be masked.

Precision check: picking the right code

Method. To test precision, we drew two random samples of 50 articles each from March 2022: one from the 3,315 articles tagged N/PET, one from the 492 articles tagged N/CMKT (seed 42). We read each headline and judged whether the article is actually about the crude oil commodity market.

N/PET (**“Crude Oil & Petroleum Products”**) is the broad subject umbrella — it fires on any article touching an oil company. The 50-article sample surfaces substantial noise:

Clearly unrelated: “Stellantis Battery Plant Investment”, “Income Investing: Global Dividends Surge” (Barron’s), “Heard on the Street: Indian Economy on the Clock”, “Australian Budget... Balancing Act”.
Oil-company corporate news (not the oil market): “Warren Buffett Gets His Deal Mojo Back”, “Carl Icahn Reduces Occidental Stake”, “Moody’s Upgrades Murphy Oil”, “ConocoPhillips to Hold Earnings Call”, “VP Lannie Sells 40,800 Of APA Corp”, “Equinor ASA: Share buy-back”, “Exxon Mobil Down Over 6%”.

N/CMKT (**“Crude Spot Market Commentary”**) is the **right code**. In the 50-article sample: the entire list is crude-market commentary — WTI/Brent price moves, OPEC decisions, inventory draws, SPR releases, Russian oil sanctions, tanker charters. **Zero** clearly-unrelated articles; zero corporate filings or earnings-call press releases.

Trade-off: N/CMKT is much narrower (~15% of N/PET’s volume). Higher precision, lower recall. The recall story on the next slides applies either way.

DJNW Metadata Audit

└ Motivation: DJNW metadata is unreliable

└ Precision check: picking the right code

Full N/PET sample: results/npet_precision_sample.txt.

Full N/CMKT sample: results/ncmkt_precision_sample.txt.

Both: 50 random articles from 2022-03, seed 42.

Populations: N/PET = 3,315 articles, N/CMKT = 492 articles.

2026-04-17

Method. To test precision, we drew two random samples of 50 articles each from March 2022: one from the 3,315 articles tagged N/PET, one from the 492 articles tagged N/CMKT (seed 42). We read each headline and judged whether the article is actually about the crude oil commodity market.

N/PET (“Crude Oil & Petroleum Products”) is the broad subject umbrella — it fires on any article touching an oil company. The 50-article sample surfaces substantial noise.

Clearly unrelated: “Stellantis Battery Plant Investment”, “Income Investing: Global Dividends Surge” (Barron’s), “Heard on the Street: Indian Economy on the Clock”, “Australian Budget... Balancing Act”.
Oil-company corporate news (not the oil market): “Warren Buffett Gets His Deal Mojo Back”, “Carl Icahn Reduces Occidental Stake”, “Moody’s Upgrades Murphy Oil”, “ConocoPhillips to Hold Earnings Call”, “VP Lannie Sells 40,800 Of APA Corp”, “Equinor ASA: Share buy-back”, “Exxon Mobil Down Over 6%”.

N/CMKT (“Crude Spot Market Commentary”) is the right code. In the 50-article sample: the entire list is crude-market commentary — WTI/Brent price moves, OPEC decisions, inventory draws, SPR releases, Russian oil sanctions, tanker charters. Zero clearly-unrelated articles; zero corporate filings or earnings-call press releases.

Trade-off: N/CMKT is much narrower (~15% of N/PET’s volume). Higher precision, lower recall. The recall story on the next slides applies either way.

- 1. **Code choice matters.** N/PET (broad petroleum umbrella) fires on corporate news, earnings, battery-plant announcements. N/CMKT (crude market commentary) is the right filter for macro research — near-zero corporate noise — but it covers only ~15% of the articles that N/PET does.
- 2. **Recall is the binding problem.** Many articles directly about oil events carry no content code at all — Saudi oil policy, Abu Dhabi sovereign wealth losses, Exxon’s Russia exit, and every cross-asset article reacting to oil moves. Even N/PET (the broader code) caps at ~8% share during the biggest oil moves.
- 3. **Wire codes add recall with a trade-off.** Post-1996, N/IPR, N/NRG, N/D0I catch more oil-relevant articles but at the cost of precision.
- 4. **Pre-1997 data is unusable** for tag-based relevance either way: wire codes spammed, content codes sparse.
- 5. **Implication.** Researchers have to pick between a narrow, precise filter (N/CMKT) that misses most oil articles and a broad filter (N/PET) with substantial noise. Neither is a comprehensive label for macro research. An LLM pipeline that reads article text can fill the recall gap without inheriting the precision failures.

2026-04-17

DJNW Metadata Audit

└ Motivation: DJNW metadata is unreliable

└ What the crude oil exercise tells us

- 1. **Code choice matters.** N/PET (broad petroleum umbrella) fires on corporate news, earnings, battery-plant announcements. N/CMKT (crude market commentary) is the right filter for macro research — near-zero corporate noise — but it covers only ~15% of the articles that N/PET does.
- 2. **Recall is the binding problem.** Many articles directly about oil events carry no content code at all — Saudi oil policy, Abu Dhabi sovereign wealth losses, Exxon’s Russia exit, and every cross-asset article reacting to oil moves. Even N/PET (the broader code) caps at ~8% share during the biggest oil moves.
- 3. **Wire codes add recall with a trade-off.** Post-1996, N/IPR, N/NRG, N/D0I catch more oil-relevant articles but at the cost of precision.
- 4. **Pre-1997 data is unusable** for tag-based relevance either way: wire codes spammed, content codes sparse.
- 5. **Implication.** Researchers have to pick between a narrow, precise filter (N/CMKT) that misses most oil articles and a broad filter (N/PET) with substantial noise. Neither is a comprehensive label for macro research. An LLM pipeline that reads article text can fill the recall gap without inheriting the precision failures.

DJNW tagging scheme: full census

2026-04-17

DJNW Metadata Audit
└ DJNW tagging scheme: full census

DJNW tagging scheme: full
census

What tags does DJNW actually put on articles?

We ran a full census across all 31.7M articles (1979–2025) to catalog every unique code that appears in the data, cross-referenced against the documentation CSVs provided by Dow Jones.

Articles carry tags in **five taxonomy fields**:

Field	Prefix	Purpose	Example
subject	N/	Topic / content / wire	N/PET, N/OPC
industry	I/	Industry sector	I/OIL, I/DRG
market	M/	Market classification	M/ENE, M/NND
product	P/	Product/feed type	P/AEQI, P/HDL
geo	R/	Geographic region	R/US, R/EU

Plus two security-identifier fields (company, isin) that carry ticker symbols and ISIN codes — not part of the DJ taxonomy.

Three documented code families (S/ statistical, G/ government, J/ journal) exist in the CSVs but **never appear on any article** in the 46-year archive.

DJNW Metadata Audit

- DJNW tagging scheme: full census

- What tags does DJNW actually put on articles?

What tags does DJNW actually put on articles?

We ran a full census across all 31.7M articles (1979–2025) to catalog every unique code that appears in the data, cross-referenced against the documentation CSVs provided by Dow Jones.

Field	Prefix	Purpose	Example
subject	N/	Topic / content / wire	N/PET, N/OPC
industry	I/	Industry sector	I/OIL, I/DRG
market	M/	Market classification	M/ENE, M/NND
product	P/	Product/feed type	P/AEQI, P/HDL
geo	R/	Geographic region	R/US, R/EU

Plus two security-identifier fields (company, isin) that carry ticker symbols and ISIN codes — not part of the DJ taxonomy.

Three documented code families (S/ statistical, G/ government, J/ journal) exist in the CSVs but **never appear on any article** in the 46-year archive.

Documentation coverage is poor

Of the unique codes that actually appear on articles, how many are documented in the Dow Jones CSVs?

Field	Unique codes	Documented	Undocumented	Coverage
subject	3,315	644	2,671	19.4%
industry	1,047	254	793	24.3%
geo	1,101	341	760	31.0%
market	368	177	191	48.1%
product	683	402	281	58.9%
All taxonomy	6,514	1,818	4,696	27.9%

Only 28% of the taxonomy codes used on articles are documented. Subject codes are the worst at 19% — and subject is the primary field for content classification.

2026-04-17

DJNW Metadata Audit

└ DJNW tagging scheme: full census

└ Documentation coverage is poor

Source: `results/djnw_tag_census.csv` — full per-code inventory with occurrence counts, field, documented/undocumented flag, first/last month. `results/djnw_tag_coverage_monthly.csv` — monthly time series. Generated by `scripts/djnw_tag_census.py`, job 8442488.

Documentation coverage is poor

Of the unique codes that actually appear on articles, how many are documented in the Dow Jones CSVs?

Field	Unique codes	Documented	Undocumented	Coverage
subject	3,315	644	2,671	19.4%
industry	1,047	254	793	24.3%
geo	1,101	341	760	31.0%
market	368	177	191	48.1%
product	683	402	281	58.9%
All taxonomy	6,514	1,818	4,696	27.9%

Only 28% of the taxonomy codes used on articles are documented. Subject codes are the worst at 19% — and subject is the primary field for content classification.

Top undocumented codes are massive

The most-used undocumented codes have **millions** of occurrences:

Code	Field	Occurrences	Active period
N/DJWB	subject	6,457,184	2000-04 to 2013-04
N/DJIN	subject	5,817,985	2001-11 to 2013-04
P/AS04	product	3,938,682	1979-06 to 2003-09
N/DJM0	subject	3,021,006	2002-08 to 2013-05
P/AS03	product	2,859,167	1996-11 to 2002-02
N/DJFP	subject	2,224,590	2004-03 to 2013-04
P/EQA	product	1,815,820	2019-04 to 2025-08
N/DJCB	subject	1,784,245	2002-06 to 2013-05
M/FIN	market	706,190	1997-01 to 2006-02
M/BTC	market	2,147	2022-07 to 2025-08
M/ETH	market	1,442	2022-07 to 2025-08

M/BTC and M/ETH (Bitcoin, Ethereum) are macro-relevant market codes introduced in 2022 — undocumented. M/FIN (presumably “Financial”) was used on 706k articles over 9 years — also undocumented.

DJNW Metadata Audit

└─ DJNW tagging scheme: full census

└─ Top undocumented codes are massive

2026-04-17

Top undocumented codes are massive			
The most-used undocumented codes have millions of occurrences			
Code	Field	Occurrences	Active period
N/DJWB	subject	6,457,184	2000-04 to 2013-04
N/DJIN	subject	5,817,985	2001-11 to 2013-04
P/AS04	product	3,938,682	1979-06 to 2003-09
N/DJM0	subject	3,021,006	2002-08 to 2013-05
P/AS03	product	2,859,167	1996-11 to 2002-02
N/DJFP	subject	2,224,590	2004-03 to 2013-04
P/EQA	product	1,815,820	2019-04 to 2025-08
N/DJCB	subject	1,784,245	2002-06 to 2013-05
M/FIN	market	706,190	1997-01 to 2006-02
M/BTC	market	2,147	2022-07 to 2025-08
M/ETH	market	1,442	2022-07 to 2025-08
M/BTC and M/ETH (Bitcoin, Ethereum) are macro-relevant market codes introduced in 2022 — undocumented. M/FIN (presumably “Financial”) was used on 706k articles over 9 years — also undocumented.			

Codes appear and disappear over time

Many undocumented codes have bounded lifespans — they were introduced, used for years, then retired without any documentation trail:

- P/AS04: 3.9M occurrences, 1979–2003, then gone
- P/AS03: 2.9M occurrences, 1996–2002 only
- N/DJWB: 6.5M occurrences, 2000–2013, then gone
- P/EQA: 1.8M occurrences, 2019–present (new, undocumented)
- P/LLM: 649k occurrences, 2023–present (new, undocumented)

The tagging scheme is a **living system** with codes being added and retired over the 46-year archive. The documentation CSVs are a **snapshot** (likely circa 2014–2017 based on coverage patterns) that does not track these changes.

Minor data quality issues: a few codes appear in the wrong field; the two notable ones are P/AGF (187k occurrences in subject) and I/UKSC (125k occurrences in subject).

2026-04-17

DJNW Metadata Audit

└ DJNW tagging scheme: full census

└ Codes appear and disappear over time

Codes appear and disappear over time

Many undocumented codes have bounded lifespans — they were introduced, used for years, then retired without any documentation trail:

- P/AS04: 3.9M occurrences, 1979–2003, then gone
- P/AS03: 2.9M occurrences, 1996–2002 only
- N/DJWB: 6.5M occurrences, 2000–2013, then gone
- P/EQA: 1.8M occurrences, 2019–present (new, undocumented)
- P/LLM: 649k occurrences, 2023–present (new, undocumented)

The tagging scheme is a **living system** with codes being added and retired over the 46-year archive. The documentation CSVs are a **snapshot** (likely circa 2014–2017 based on coverage patterns) that does not track these changes.

Minor data quality issues: a few codes appear in the wrong field; the two notable ones are P/AGF (187k occurrences in subject) and I/UKSC (125k occurrences in subject).

Implications for the pipeline

- 1. **DJNW tags are not a reliable ground truth** for article–asset relevance. Any evaluation framework that treats DJ metadata as labels will inherit systematic gaps.
- 2. **The N/ namespace is especially problematic:** it mixes content tags, wire codes, and undocumented codes, with only 19% documentation coverage.
- 3. **New macro-relevant codes are being added without documentation** (M/BTC, M/ETH, P/LLM). The docs CSVs are stale.
- 4. **Our LLM pipeline reads article text**, not metadata — it is immune to these tagging gaps. This is a structural advantage over any approach that relies on DJNW codes for filtering or labeling.

2026-04-17

- DJNW Metadata Audit
 - └ DJNW tagging scheme: full census
 - └ Implications for the pipeline

Implications for the pipeline	
1.	DJNW tags are not a reliable ground truth for article–asset relevance. Any evaluation framework that treats DJ metadata as labels will inherit systematic gaps.
2.	The N/ namespace is especially problematic: it mixes content tags, wire codes, and undocumented codes, with only 19% documentation coverage.
3.	New macro-relevant codes are being added without documentation (M/BTC, M/ETH, P/LLM). The docs CSVs are stale.
4.	Our LLM pipeline reads article text , not metadata — it is immune to these tagging gaps. This is a structural advantage over any approach that relies on DJNW codes for filtering or labeling.